

Deep Learning and Convolutional Neural Network (42028)

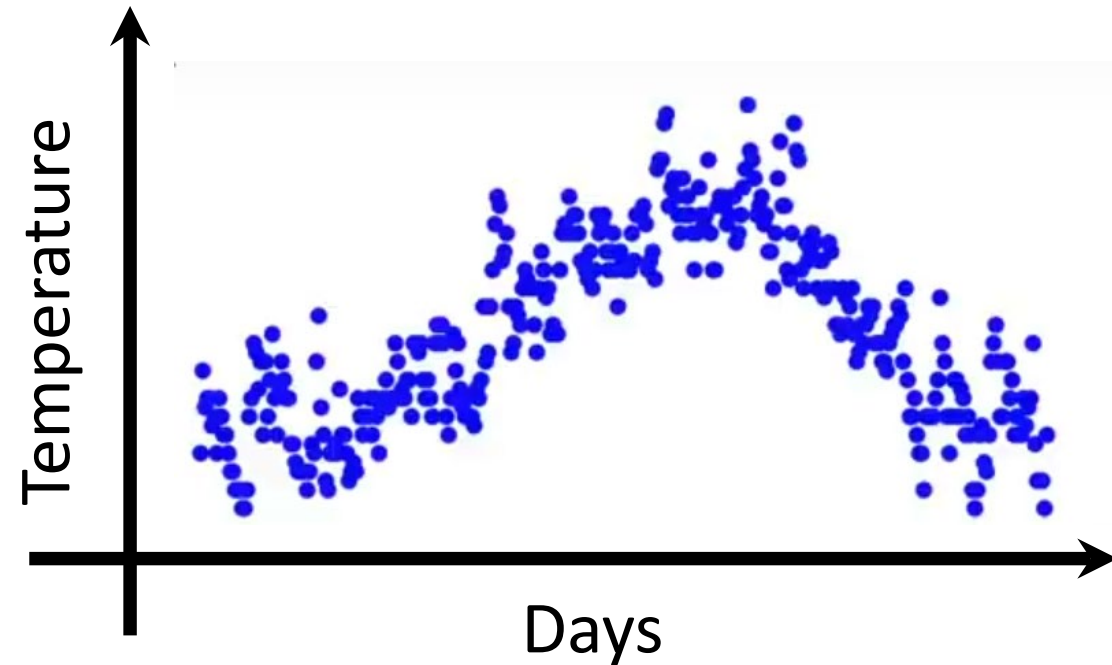
Convolutional Neural Network (CNN) - 3

Optimizers

- SGD with momentum
- RMSProp
- Adam

Exponentially Weighted Averages

- One of the popular algorithm for smoothing sequential data
- Also called Moving Average
- Weight the number of observations and using their average
- Example:
Temperature over θ days



Exponentially Weighted Averages

V_t : Moving average on day 't'

So, let

$$V_0 = 0$$

$$V_1 = 0.9 V_0 + 0.1 \theta_1$$


$$V_2 = 0.9 V_1 + 0.1 \theta_2$$

$$V_3 = 0.9 V_2 + 0.1 \theta_3$$

:

:

$$V_t = 0.9 V_{t-1} + 0.1 \theta_t$$




Exponentially Weighted Averages

$$V_t = 0.9 V_{t-1} + 0.1 \theta_t$$

If $\beta = 0.9$,

$$V_t = \beta V_{t-1} + (1 - \beta) \theta_t$$

This equation gives the moving average shown by the red line.



Exponentially Weighted Averages

$$V_t = \beta V_{t-1} + (1 - \beta) \theta_t$$

V_t is approximate average over


$\approx \frac{1}{1 - \beta}$ days

So,

$\beta = 0.9$ is closer to 10 days temperature

$\beta = 0.98$ is closer to 50 days temperature

$\beta = 0.5$ is closer to 2 days temperature



Exponentially Weighted Averages

What is Exponentially Weighted Averages doing?

$$V_t = \beta V_{t-1} + (1-\beta) \theta_t$$

For,

$$V_{100} = 0.9 V_{99} + 0.1 \theta_{100}$$

$$V_{99} = 0.9 V_{98} + 0.1 \theta_{99}$$

Substituting, V_{99}

$$V_{100} = 0.1 \theta_{100} + 0.9 (0.9 V_{98} + 0.1 \theta_{99})$$

$$V_{100} = 0.1 \theta_{100} + 0.9 (0.1 \theta_{99} + 0.9 (0.9 V_{97} + 0.1 V_{98})) ..$$

Optimizers – SGD with Momentum

- *“Compute the Exponentially weighted average of the gradients and use that gradient to update weights” - **Andrew NG***
- One of the most popular algorithms
- Helps to accelerate the gradient vectors in right direction and reduces oscillation
- Always faster than the SGD

Optimizers – SGD with Momentum

Algorithm:

At iteration t :

Calculate dw and db on the current mini-batch

$$V_{dw} = \beta V_{dw} + (1 - \beta) dw \rightarrow V_t = \beta V_{t-1} + (1 - \beta) \theta_t$$

$$V_{db} = \beta V_{db} + (1 - \beta) db$$

Update w and b :


$$w = w - \alpha V_{dw}, b = b - \alpha V_{db}$$

Hyper-parameters: α, β

Optimizers – SGD with Momentum



SGD Without Momentum



SGD With Momentum

Faster convergence and reduced oscillation

Optimizers – RMSProp


- Root Mean Square Propagation
- Unpublished adaptive learning method by Geoffery Hinton
- RMSProp also reduces oscillation but in a different way than Momentum
- RMSprop as well divides the learning rate by an exponentially decaying average of squared gradients.

Optimizers – RMSProp

Algorithm:


At iteration t:

Calculate dw and db on the current mini-batch

$$S_{dw} = \beta_2 S_{dw} + (1 - \beta_2) dw^2$$


$S_{db} = \beta_2 S_{db} + (1 - \beta_2) db^2$

Update w and b :

$$w = w - \alpha \frac{dw}{\sqrt{S_{dw}}}, \quad b = b - \alpha \frac{db}{\sqrt{S_{db}}}$$


Optimizers – RMSProp

Intuition:

$S_{dw} \rightarrow$ Smaller number expected

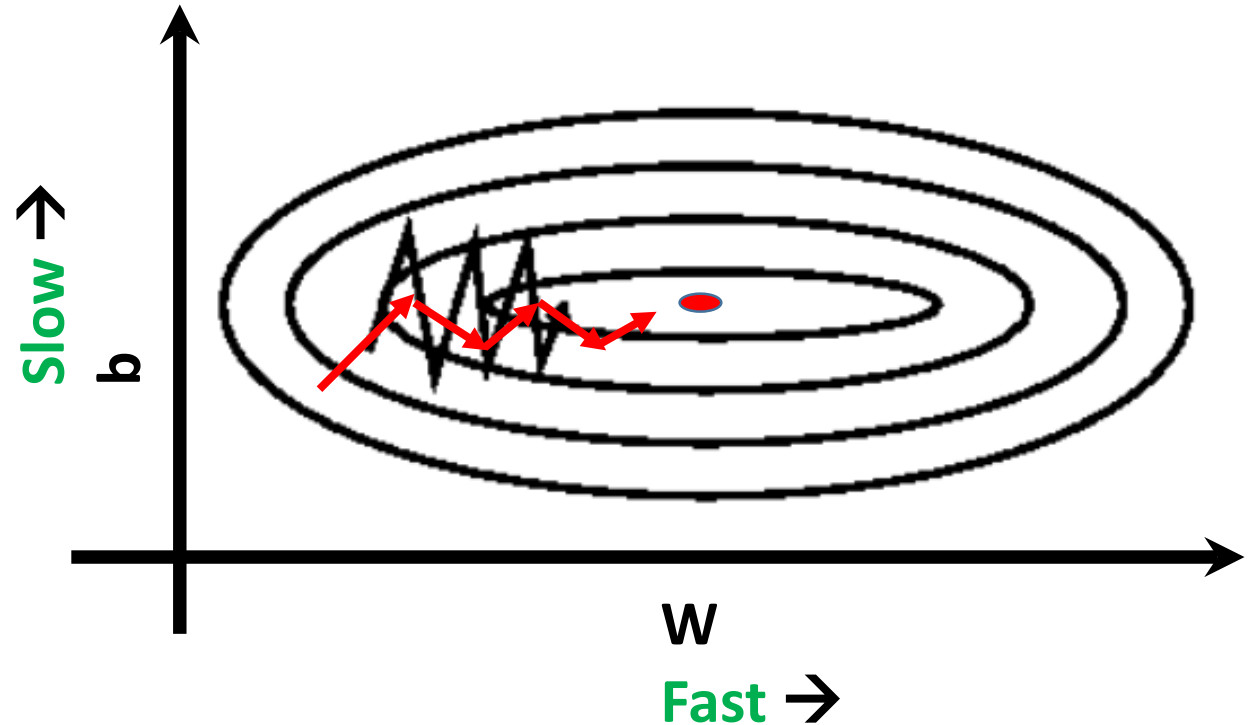
$S_{db} \rightarrow$ Larger number expected

So,

$$w = w - \alpha \frac{dw}{\sqrt{S_{dw}}}, \quad b = b - \alpha \frac{db}{\sqrt{S_{db}}}$$

Smaller number
So, w is larger

Larger number
So, b is small



In Practice add ϵ :

$$w = w - \alpha \frac{dw}{\sqrt{S_{dw} + \epsilon}}, \quad b = b - \alpha \frac{db}{\sqrt{S_{db} + \epsilon}}$$

$\epsilon \rightarrow$ small number, 10^{-8}

Optimizers – Adam

- Adam → Adaptive Moment Estimation
- Combination of RMSProp and Momentum
- Work well for a wide range of deep learning architecture

Optimizers – Adam

Algorithm:

Initialize $V_{dw} = 0$, $V_{db} = 0$, $S_{dw} = 0$, $S_{db} = 0$

At iteration t:

Calculate dw and db on the current mini-batch

$$V_{dw} = \beta_1 V_{dw} + (1 - \beta_1) dw, \quad V_{db} = \beta_1 V_{db} + (1 - \beta_1) db \leftarrow \text{From Momentum, } \beta_1$$

$$S_{dw} = \beta_2 S_{dw} + (1 - \beta_2) dw^2, \quad S_{db} = \beta_2 S_{db} + (1 - \beta_2) db^2 \leftarrow \text{From RMSProp, } \beta_2$$

Update w and b :

$$w = w - \alpha \frac{V_{dw}}{\sqrt{S_{dw}} + \epsilon}, \quad b = b - \alpha \frac{V_{db}}{\sqrt{S_{db}} + \epsilon}$$

Optimizers – Adam

In practice: Bias correction is required as V_{dw} , V_{db} , S_{dw} , S_{db} are initialized to 0 and are biased towards zero. Hence, a bias correction is required as follows:

$$V'_{dw} = \frac{V_{dw}}{(1 - \beta_1)}, \quad V'_{db} = \frac{V_{db}}{(1 - \beta_1)}$$
$$S'_{dw} = \frac{S_{dw}}{(1 - \beta_2)}, \quad S'_{db} = \frac{S_{db}}{(1 - \beta_2)}$$

Update w and b:

$$w = w - \alpha \frac{V'_{dw}}{\sqrt{S'_{dw}} + \epsilon}, \quad b = b - \alpha \frac{V'_{db}}{\sqrt{S'_{db}} + \epsilon}$$

Optimizers – Adam

Hyper parameter guide:

α (Learning rate) \rightarrow should be tuned, start with 0.001

β_1 (Momentum term) $\rightarrow 0.9$ (dw)

β_2 (moving weighted average) $\rightarrow 0.999$ (dw²)

$\epsilon \rightarrow 10^{-8}$

Optimization Demo:

<https://vis.ensmallen.org/>


ImageNet Dataset:

- 15+ million labelled high-resolution images
- 22000 categories
- ILSVRC (**Large Scale Visual Recognition Challenge**) used a subset of ImageNet:
 - ~1000 images per category
 - 1000 categories
 - Train: 1.2 million images
 - Validation: 50k images
 - Test : 150k images

ImageNet Dataset:



ImageNet Dataset Results:



Transfer Learning

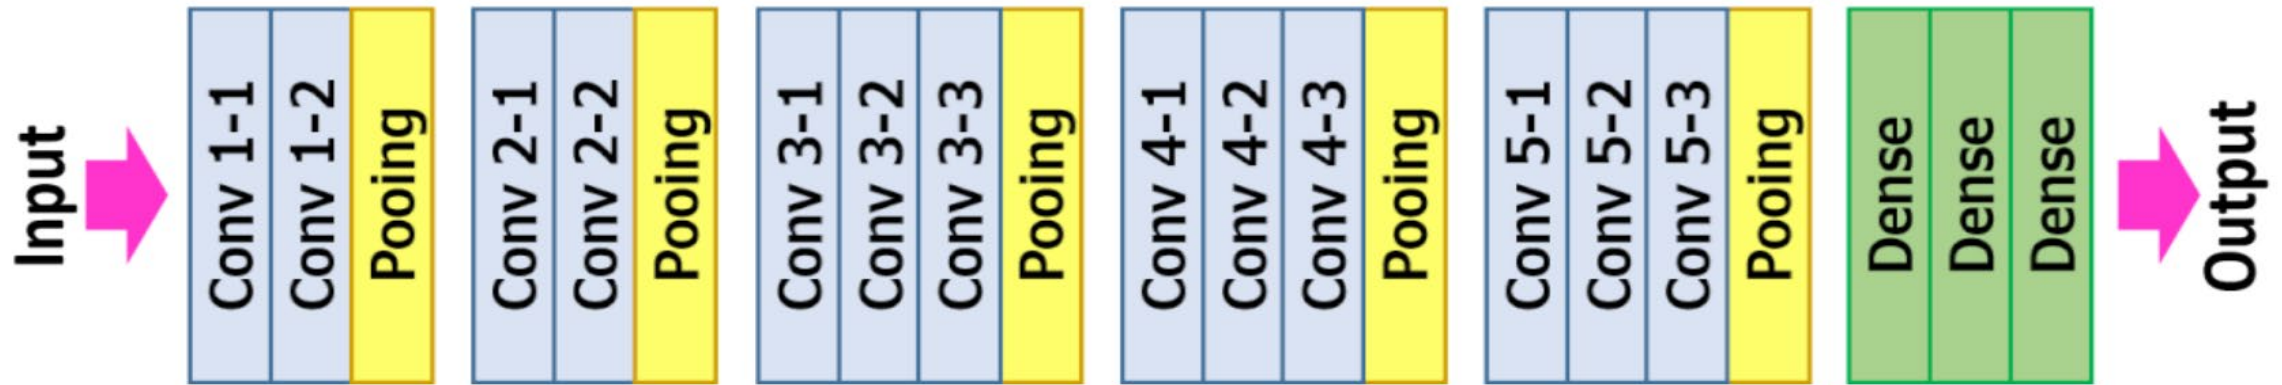
- Knowledge acquired while solving one task, can be used to solve related tasks.
- Example:
 - You know how to ride a Bi-cycle → You can learn how to ride a Motorbike
 - You know how to use a Tablet → You can easily learn how to use a Laptop/desktop
- Similar to the way humans apply knowledge acquired from one task to solve a new but similar/related task.
- We learned how to read in Year-1 in literacy class. Reading skills acquired in the literacy classes made it easy to understand Physics in Year-9.

Transfer Learning Benefits

- 1. Less training data required:** Don't have enough data to train a Deep Learning model from scratch. Model trained using a large (similar) dataset can be used.
- 2. Faster training :** Training can converge faster, due the use to existing knowledge (weights) to start with rather than from scratch.
- 3. Better model generalization:** Model is trained to identify features which can be applied to new contexts.

Transfer Learning Strategies

VGG-16

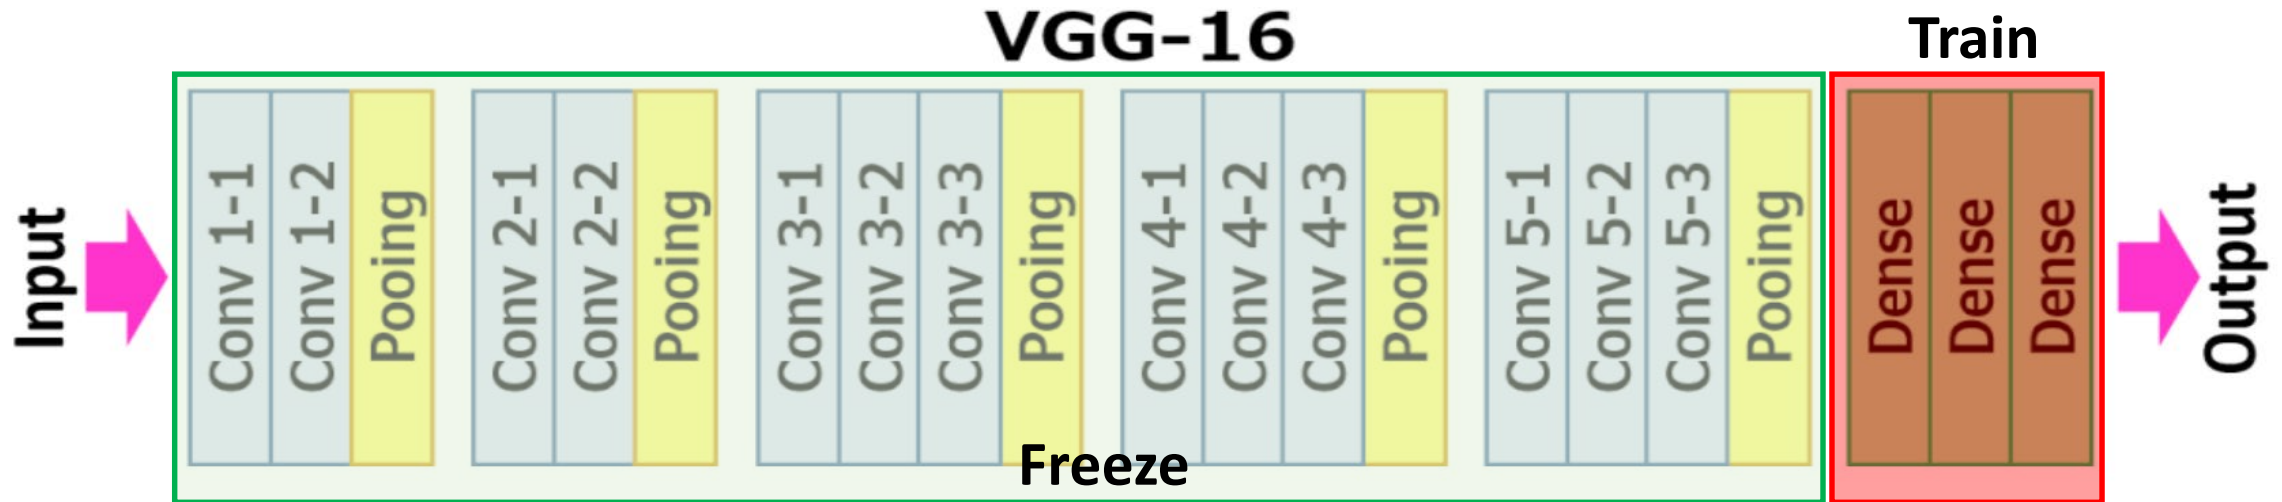


Option-1: (VGG-16 considered as an example)

Use pre-trained (ImageNet) model for prediction, without any training.

→ Useful when your dataset distribution is similar to ImageNet, with small number of samples.

Transfer Learning Strategies

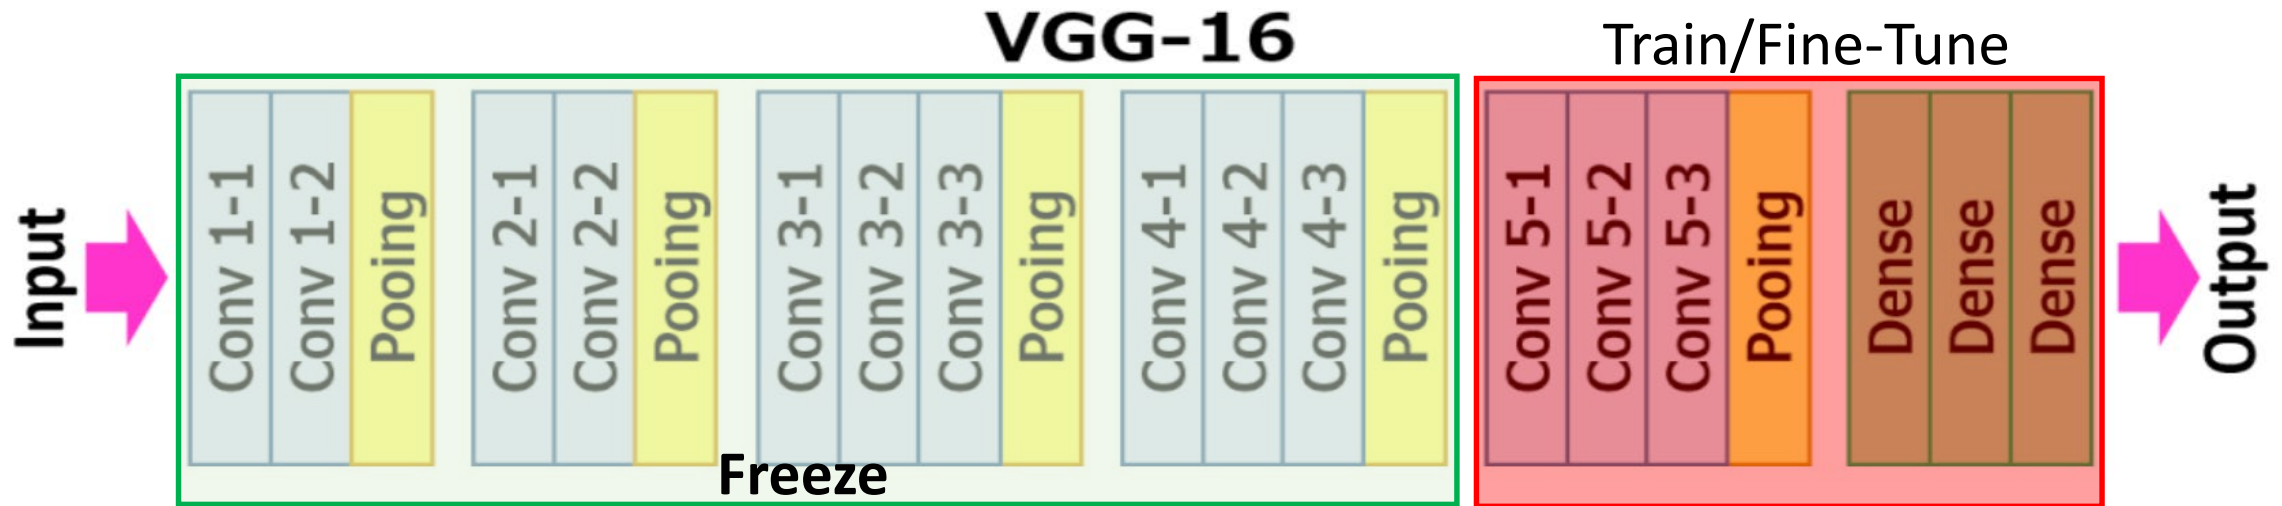


Option-2: (VGG-16 considered as an example)

Train Full-Connected layer, Use CONV layers for feature extraction

→ Useful when your dataset distribution is similar to ImageNet (or original dataset), but number of classes are different and your dataset is small.

Transfer Learning Strategies



Option-3: (VGG-16 considered as an example)

Partially Train CONV layers (usually last layer(s) which have specialised features) + Full Connection (FC) layer (with modifications)

→ Useful when your dataset distribution is not similar to ImageNet (or original dataset), number of classes are different and your dataset is small.

Transfer Learning Strategies

VGG-16



Option-4: (VGG-16 considered as an example)

Train all the CONV layers + Full Connection (FC) layer (with modifications)


→ Useful when your dataset distribution is not similar to ImageNet, number of classes are different, your dataset is large and the task is complex.

Classic CNN Architectures

Case Study: AlexNet

- Similar architecture as LeNet by Yann LeCunn et al. but deeper with more layers
- Simple architecture:
 - CONV : 5 layers
 - FC: 3 layer
 - Max pooling
 - Dropout
- Accuracy: top-5 test error rate of 15.3%
- Winner of ILSVRC 2012!
- First CNN to be successful on a very big dataset!

Case Study: AlexNet



Case Study: AlexNet

Input: 224x224x3 image

CONV1 → CONV2 → CONV3 → CONV4 → CONV5 → FC1 → FC2 → FC3

Filters: 96
Dim: 11x11
Stride: 4
Pad: 0

Filters: 256
Dim: 5x5
Stride: 1
Pad: 2

Filters: 384
Dim: 3x3
Stride: 1
Pad: 1

Filters: 384
Dim: 3x3
Stride: 1
Pad: 1

Filters: 256
Dim: 3x3
Stride: 1
Pad: 1

4096
Neuron

4096
Neuron

1000
Neuron

Activations: Relu after each CONV and FC layer

Optimizer: SGD with Momentum

Regularization: Dropout in FC1 and FC2

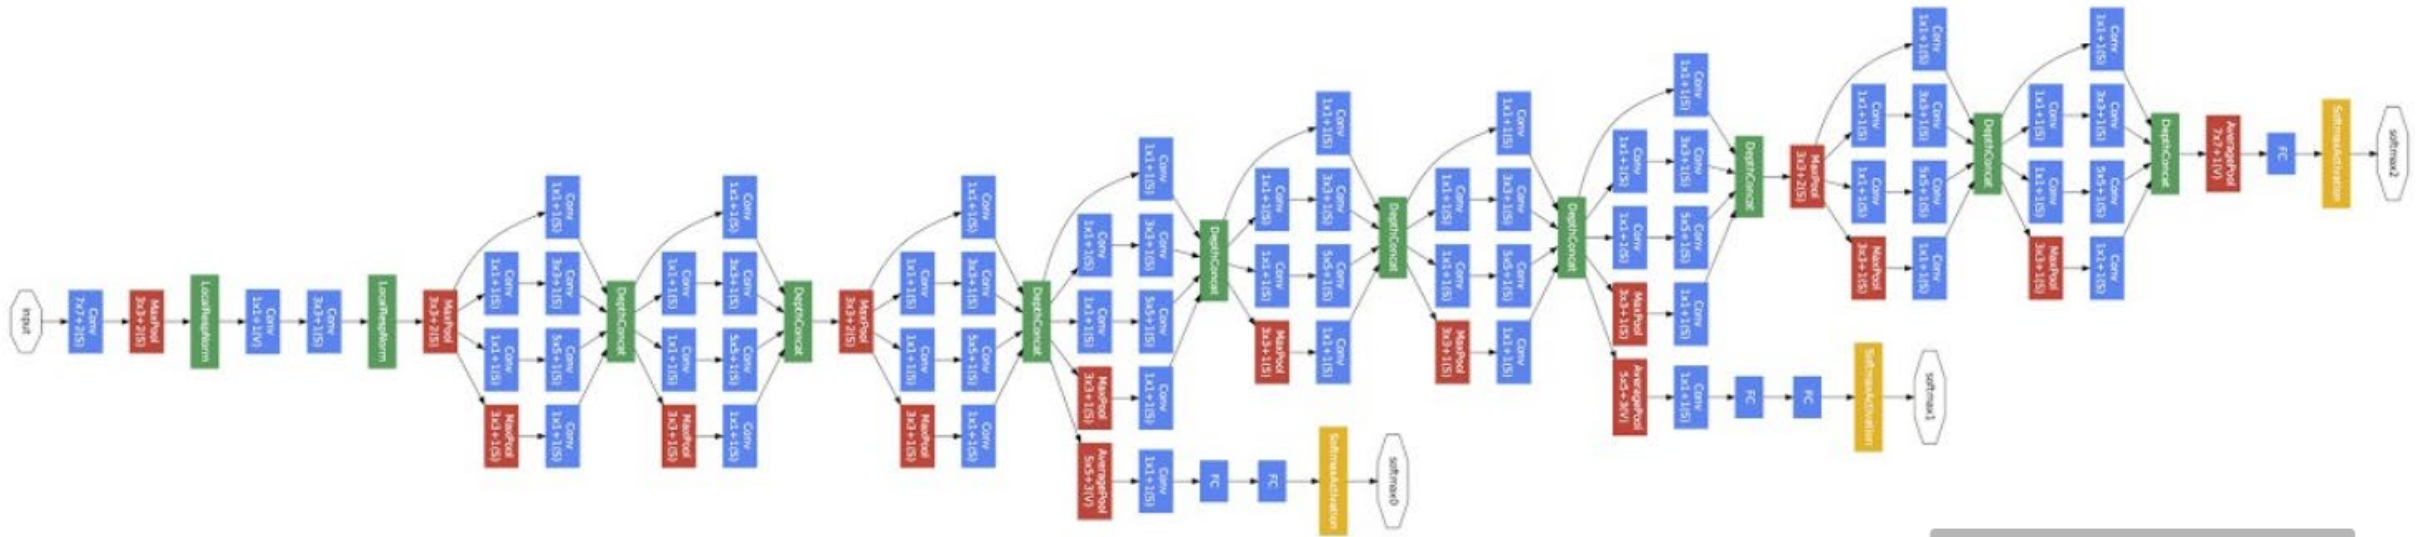
Total Trainable parameter: ~60Million

Training settings: 2 X Nvidia GTX 580 3GB GPUs for 5-6days!

Case Study: GoogleNet/Inception(2014)

- Accuracy: top-5 test error rate of 6.7%
- Close to human level performance
- Winner of ILSVRC 2014!
- 22 layer Deep CNN
- Number of trainable parameters: 4 Million (Alexnet ~ 60M), Significantly reduced
- A novel inception module was introduced.
- Optimizer: RMSProp

Case Study: GoogleNet/Inception(2014)



Convolution
Pooling
Softmax
Other

Case Study: GoogleNet/Inception(2014)

Inception Module

